

# Adaptive Quantum Error Correction

Kollapudi Jithin    Krunal Vaghela    Yaswanth R Kumar

November 23, 2025

Let's Start with a challenge...

# Coin Tossing: An Optimal Strategy Challenge

Defining the Constraint and the Goal

## The Setup

We are given three independent coins,  $C_1$ ,  $C_2$ , and  $C_3$ , each producing a binary outcome  $\{0, 1\}$ .

- **Bias:** Each coin  $C_i$  has a fixed but **unknown** probability  $p_i$  of yielding a 1.

$$P(X = 1 \mid C_i) = p_i$$

- **Budget:** We are limited to a total of  $\mathbf{N = 100}$  tosses across all three coins.

## The Question

What is the **most optimal strategy** (a sequential rule for choosing the next coin to toss) that maximizes the total expected number of **1s** observed within the 100-toss budget?

# Context: The Multi-Armed Bandit (MAB) Problem

## The Exploration-Exploitation Trade-off

This scenario is a classic example of a **Stochastic Multi-Armed Bandit (MAB)** problem. The key to the optimal solution lies in balancing two conflicting actions:

### Exploration (Learning)

- Action: Toss all coins enough times to estimate their true probabilities ( $\hat{p}_i$ ) accurately.
- Risk: Wasting valuable tosses on a coin that might be sub-optimal.

### Exploitation (Earning)

- Action: Focus the majority of the budget on the coin currently appearing to be the best ( $\hat{p}_{\max}$ ).
- Risk: Missing out on a truly optimal coin whose potential was underestimated early on.

# What is the Optimal Strategy then?

# The Goal of Optimal Strategy: Minimizing Regret

## Optimal Strategy

The solution is a sequential strategy (e.g., **UCB** or **Thompson Sampling**) designed to minimize the total expected **Regret** over the 100 tosses.

## Definition (Total Expected Regret ( $R_N$ ))

Regret is the difference between the total reward we **would have** received if we knew the best coin and the total reward received by following our chosen strategy.

Mathematically, for  $N = 100$  tosses:

$$R_{100} = \sum_{t=1}^{100} \left[ \max_i(p_i) - p_{c^{(t)}} \right]$$

where  $p_{c^{(t)}}$  is the true probability of success of the coin chosen at time  $t$ .

**Goal:** A successful MAB strategy seeks to minimize this accumulated loss over time.

# UCB Algorithm: Optimistic Exploration

The UCB algorithm is an elegant strategy that prioritizes coins (arms) that have a high estimated reward **or** have been played too few times (high uncertainty). It systematically cycles through exploration and exploitation based on calculated confidence intervals.

## The UCB Selection Rule

At each step  $t$ , the coin  $C_i$  to toss next is the one that maximizes its **UCB value**  $A_i(t)$ :

$$\text{Choose } C_i = \arg \max_i A_i(t)$$

where  $A_i(t)$  is calculated as:

$$A_i(t) = \underbrace{\hat{p}_i}_{\text{Exploitation Term}} + \underbrace{\sqrt{\frac{2 \ln t}{n_i(t)}}}_{\text{Exploration Term}}$$

$$A_i(t) = \underbrace{\hat{p}_i}_{\text{Exploitation Term}} + \underbrace{\sqrt{\frac{2 \ln t}{n_i(t)}}}_{\text{Exploration Term}}$$

- $\hat{p}_i$ : The current average reward (estimated probability of getting a 1) for coin  $C_i$ . (*The best estimate so far.*)
- $n_i(t)$ : The number of times coin  $C_i$  has been tossed up to time  $t$ .
- $\ln t$ : The natural logarithm of the total number of tosses made so far.

## How it Balances Trade-off

- **Exploitation:** By using  $\hat{p}_i$ , it favors coins with high success rates.
- **Exploration:** The exploration term decreases as  $n_i(t)$  increases. Coins that haven't been played much (low  $n_i(t)$ ) get a large boost, forcing the algorithm to be "optimistic" about their potential and try them out.

Why do all this? Why not just simple toss each coin a few times at the beginning and pick the one with most successes?

Let's see a few numbers

**Table:** Total Cumulative Regret for MAB Strategies

<b>Testcase</b>	<b><math>\epsilon</math>-Greedy Regret</b>	<b>UCB Regret</b>	<b>KL-UCB Regret</b>	<b>Thompson Regret</b>
1	55.10	24.73	7.32	4.78
2	410.85	263.19	75.25	51.59
3	950.40	571.33	134.06	74.61

Averaged results from simulations for 30 coins for an horizon of 10000-250000 tosses

Back to QEC...

Classical wireless standards (Wi-Fi/5G) use **Adaptive Modulation and Coding (AMC)** to handle varying noise.

- *Low Noise:* Use high-rate schemes like **64-QAM**.
- *High Noise:* Switch to robust schemes like **QPSK** or **BPSK**.

*We aim to apply this same adaptive principle to Quantum Error Correction.*

Can we do this in Quantum Networks?

**Problem:** Transmitting quantum information over noisy channels

- Depolarizing channel with unknown noise rate  $p$
- Multiple QEC codes available with different costs
- **Goal:** Maximize fidelity while minimizing physical qubit overhead

**Key Challenge:** How do we adaptively choose the right quantum code in real-time?

# Depolarizing Channel Model

The depolarizing channel applies noise to transmitted qubits:

$$\mathcal{E}_p(\rho) = (1 - p)\rho + \frac{p}{3}(X\rho X + Y\rho Y + Z\rho Z) \quad (1)$$

- $p$  is the **physical error rate** (unknown to us)
- With probability  $(1 - p)$ : no error
- With probability  $p/3$  each: bit-flip ( $X$ ), phase-flip ( $Z$ ), or both ( $Y$ )

**Challenge:** We must learn  $p$  online while transmitting qubits

# Choosing the Quantum Error-Correcting Codes

Say we have three quantum error-correcting codes to choose from:

Code Name	Qubits $C_k$	Trade-off
5-qubit perfect	5	Lowest cost, less robust
7-qubit Steane	7	Medium cost/protection
9-qubit Shor	9	Highest cost, most robust

These codes were selected because they all share a crucial property, a minimum distance of  $d = 3$ . This distance guarantees that each code can correct **at least one arbitrary single-qubit Pauli error** (X, Y, or Z).

**Key Insight:** The "best" code depends on the unknown noise level  $p$

- Low noise → cheaper codes suffice
- High noise → need expensive protection

# Multi-Armed Bandit Formulation

## Bandit Framework:

- **Arms:**  $k \in \{5, 7, 9\}$  (the three QEC codes)
- **Time horizon:**  $T$  logical qubits to transmit
- **Action at time  $t$ :** Choose code  $k_t$
- **Observation:** Logical error indicator  $\ell_t \in \{0, 1\}$
- **Fidelity:**  $f_t = 1 - \ell_t$  (success = 1, failure = 0)

**Expected fidelity** for code  $k$  under channel  $p$ :

$$F_k(p) = 1 - L_k(p) \tag{2}$$

where  $L_k(p) = \Pr(\text{logical error} \mid \text{code } k, \text{ noise } p)$

What is a reward in this setting?

# What is a reward in this setting?

Less number of qubits?

# What is a reward in this setting?

Less number of qubits?  
High Fidelity?

# Reward Definition: Fidelity vs Cost

We define a scalar reward trading off fidelity against qubit cost:

$$r_t = f_t - \lambda C_{k_t}$$

(3)

where:

- $f_t$ : observed fidelity (0 or 1)
- $C_{k_t}$ : physical qubits used
- $\lambda > 0$ : cost weight (how many fidelity points = one qubit)

**Expected reward** for code  $k$  under noise  $p$ :

$$\mu_k(p) = \mathbb{E}[r_t | k_t = k] = F_k(p) - \lambda C_k \quad (4)$$

Higher  $\lambda \rightarrow$  prefer cheaper codes | Lower  $\lambda \rightarrow$  prefer robust codes

# Oracle and Regret Definition

**Best fixed code in hindsight** (knows true  $p$ ):

$$k^*(p) = \arg \max_{k \in \{5,7,9\}} (F_k(p) - \lambda C_k) \quad (5)$$

**Regret:** Cumulative loss compared to oracle

$$R_T = \sum_{t=1}^T \mu_{k^*(p)}(p) - \sum_{t=1}^T r_t \quad (6)$$

Equivalently:

$$R_T = \sum_{t=1}^T [\mu_{k^*(p)}(p) - \mu_{k_t}(p)] \quad (7)$$

**Goal:** Minimize  $R_T$  (learn the best code quickly)

# Regret Decomposition

We can break regret into two interpretable components:

**Fidelity Regret:**

$$R_T^{(F)} = \sum_{t=1}^T F_{k^*(p)}(p) - \sum_{t=1}^T f_t \quad (8)$$

**Cost (Qubit) Regret:**

$$R_T^{(Q)} = \sum_{t=1}^T C_{k_t} - \sum_{t=1}^T C_{k^*(p)} \quad (9)$$

**Relationship:**

$$R_T = R_T^{(F)} - \lambda \cdot R_T^{(Q)} \quad (10)$$

# Estimating the Channel: $\hat{p}_t$

**Model-Based Approach:** Estimate  $p$  from observed logical errors

Given data up to time  $t$ :  $\mathcal{D}_t = \{(k_\tau, \ell_\tau)\}_{\tau=1}^t$

**Maximum Likelihood Estimation:**

$$\hat{p}_t = \arg \max_{p \in [0, p_{\max}]} \prod_{\tau=1}^t L_{k_\tau}(p)^{\ell_\tau} \cdot (1 - L_{k_\tau}(p))^{1-\ell_\tau} \quad (11)$$

Or equivalently, maximize log-likelihood:

$$\hat{p}_t = \arg \max_p \sum_{\tau=1}^t \left[ \ell_\tau \log L_{k_\tau}(p) + (1 - \ell_\tau) \log(1 - L_{k_\tau}(p)) \right] \quad (12)$$

**Alternative:** Least-squares fit to observed fidelities

$$\hat{p}_t = \arg \min_p \sum_{\tau=1}^t (f_\tau - F_{k_\tau}(p))^2 \quad (13)$$

# Strategy 1: Pure Exploitation (Greedy)

Given current estimate  $\hat{p}_t$ , compute estimated rewards:

$$\hat{\mu}_{k,t} = F_k(\hat{p}_t) - \lambda C_k \quad (14)$$

**Greedy policy:** Choose the code with maximum estimated reward

$$k_{t+1} = \arg \max_{k \in \{5,7,9\}} \hat{\mu}_{k,t} \quad (15)$$

## Pros:

- Simple to implement
- Exploits current knowledge optimally

## Cons:

- No explicit exploration
- Can get stuck on suboptimal code if early estimates mislead

## Strategy 2: Model-Free UCB1

**Direct bandit learning** without estimating  $p$ :

Track for each arm  $k$ :

- Number of pulls:  $N_{k,t} = \sum_{\tau=1}^t \mathbb{I}\{k_\tau = k\}$
- Empirical mean reward:  $\hat{\mu}_{k,t} = \frac{1}{N_{k,t}} \sum_{\tau: k_\tau = k} r_\tau$

**UCB1 policy:**

$$k_{t+1} = \arg \max_k \left( \hat{\mu}_{k,t} + \sqrt{\frac{2 \log t}{N_{k,t}}} \right) \quad (16)$$

- Exploration bonus  $\propto 1/\sqrt{N_{k,t}}$
- Proven  $O(\log T)$  regret bound
- Doesn't require knowing  $F_k(p)$  functional form

# Strategies : Comparison

- **Greedy (Exploitation):** This strategy is Model-Based (uses  $\hat{p}_t$ ) but includes zero explicit exploration. It risks long-term losses by relying entirely on the accuracy of the current best guess of the noise, potentially getting stuck on a suboptimal code.
- **Model-Free UCB1 (Guaranteed Exploration):** This strategy is Model-Free (ignores  $\hat{p}_t$ ), learning only from observed rewards. It uses an explicit exploration bonus based on pull count, offering superior robustness and providing provable  $\mathbf{O}(\log T)$  regret bounds.

# Connection to Our Heuristic

## Current approach: Window-based syndrome counting

- ① Count syndrome-triggered logical errors in a recent window.
- ② Form an empirical noise estimate  $\tilde{p}_t$ .
- ③ Apply fixed decision thresholds:
  - $\tilde{p}_t$  small  $\Rightarrow$  use 5-qubit code,
  - $\tilde{p}_t$  moderate  $\Rightarrow$  use 7-qubit code,
  - $\tilde{p}_t$  large  $\Rightarrow$  use 9-qubit code.

## Connection to Our Heuristic Contd.

### Mathematical interpretation:

This rule is equivalent to a **greedy parametric bandit policy** with a piecewise threshold structure:

$$k_{t+1} = \begin{cases} 5, & \tilde{p}_t < p_{(5,7)}, \\ 7, & p_{(5,7)} \leq \tilde{p}_t < p_{(7,9)}, \\ 9, & \tilde{p}_t \geq p_{(7,9)}, \end{cases}$$

where  $p_{(5,7)}$  and  $p_{(7,9)}$  are the noise values at which the expected rewards satisfy

$$\mu_5(p_{(5,7)}) = \mu_7(p_{(5,7)}), \quad \mu_7(p_{(7,9)}) = \mu_9(p_{(7,9)}),$$

with  $\mu_k(p) = F_k(p) - \lambda C_k$ .

# Interpreting the Window Size $W$

- A window size  $W=4$  **does not** mean sending classical feedback every 4 *physical* qubits.  
**It means feedback is generated once every 4 logical qubits.**
- The choice of  $W$  reflects an exploration–exploitation tradeoff: smaller  $W$  gives faster adaptation and better early estimates of  $p$ , but increases classical communication overhead.
- If total communication time is part of the objective, we can refine the regret metric by incorporating both:

$$t_Q \propto C_k \quad (\text{quantum transmission time}),$$

$$t_C \propto \frac{1}{W} \quad (\text{classical feedback rate}).$$

- To reduce overhead as the estimate of  $p$  stabilizes, we may let the window size grow with time, e.g.

$$W_t = W_0 + \alpha \ln(t+1),$$

providing rapid early exploration and slower, less costly feedback at later times.

## Interpreting the Window Size $W$ (Contd.)

Consider a time instant  $t$  where we transmit a logical state  $\psi_L^{(t)}$  encoded using the code  $C_k^{(t)}$ . A natural question is:

**Does the window size  $W$  need to be the same for all codes?**

- Higher-redundancy codes (e.g., the 9-qubit Shor code) produce fewer independent syndrome events per logical transmission. Hence they provide **less information per logical qubit** about the underlying noise rate  $p$ .
- Lower-redundancy codes (e.g., the 5-qubit perfect code) accumulate statistically useful syndrome data more rapidly.
- It is therefore reasonable to use a **code-dependent window size**

$$W_k \quad \text{with} \quad W_9 < W_7 < W_5,$$

where more redundant codes use smaller windows to compensate for their lower information rate.

- A practical choice is:  $W_9 = 1$ ,  $W_7 = 2$ ,  $W_5 = 3$ , ensuring that each code contributes approximately comparable statistical confidence in its noise estimate.

# Non-Stationary Channels: Sliding Windows

**Real-world complication:** Physical error rate  $p$  may drift over time

- Temperature fluctuations
- Hardware degradation
- External interference

**Solution:** Use **sliding window** estimation

- Only use recent  $W$  logical transmissions to estimate  $\hat{p}_t$
- Discards old data that may reflect outdated channel conditions
- Trade-off: smaller  $W$  tracks changes faster but has higher variance

**Modified MLE:**

$$\hat{p}_t = \arg \max_p \sum_{\tau=\max(1,t-W+1)}^t \left[ \ell_\tau \log L_{k_\tau}(p) + (1 - \ell_\tau) \log(1 - L_{k_\tau}(p)) \right] \quad (17)$$

# Algorithm Summary

---

**Algorithm 1** Adaptive QEC Code Selection

---

- 1: **Input:** Codes  $\{5, 7, 9\}$ , cost weight  $\lambda$ , horizon  $T$ , window  $W$
- 2: Initialize:  $\hat{p}_0 = 0.5$ ,  $\mathcal{D} = \emptyset$
- 3: **for**  $t = 1$  to  $T$  **do**
- 4:   Estimate channel:  $\hat{p}_t = \text{MLE}(\mathcal{D}_{\text{recent}})$
- 5:   Compute rewards:  $\hat{\mu}_k = F_k(\hat{p}_t) - \lambda C_k$  for each  $k$
- 6:   **Choose code:**  $k_t = \arg \max_k \hat{\mu}_k$    (or UCB variant)
- 7:   Encode logical qubit using code  $k_t$
- 8:   Transmit through depolarizing channel
- 9:   Decode and measure: observe  $\ell_t \in \{0, 1\}$
- 10:   Compute fidelity:  $f_t = 1 - \ell_t$
- 11:   Compute reward:  $r_t = f_t - \lambda C_{k_t}$
- 12:   Update dataset:  $\mathcal{D} \leftarrow \mathcal{D} \cup \{(k_t, \ell_t)\}$
- 13:   Keep only recent window:  $\mathcal{D}_{\text{recent}} \leftarrow \text{last } W \text{ samples}$
- 14: **end for**

## What can we measure:

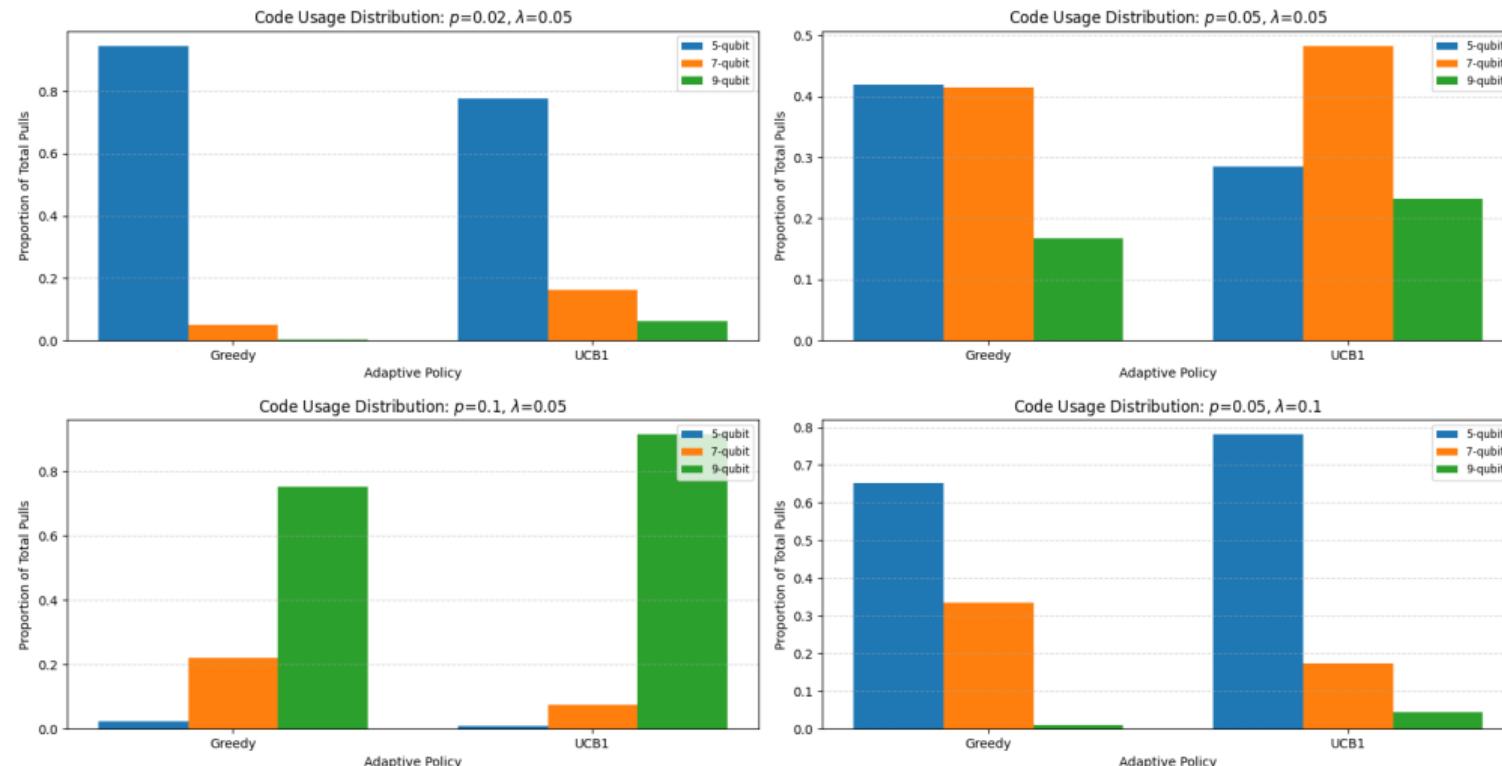
- ① **Cumulative Regret:**  $R_T = \sum_{t=1}^T (\mu_{k^*} - r_t)$ 
  - Should grow sublinearly:  $R_T = O(\log T)$  for good policies
- ② **Regret Decomposition:**  $(R_T^{(Q)}, R_T^{(F)})$
- ③ **Convergence Rate:** Time to identify  $k^*(p)$ 
  - How quickly does  $\hat{p}_t \rightarrow p$ ?

# Expected Results

**Hypothesis:** Our bandit-based approach will:

- **Outperform fixed policies** that always use one code
  - Lower regret across different noise regimes
- **Approach oracle performance** as  $T \rightarrow \infty$ 
  - $R_T/T \rightarrow 0$  (per-step regret vanishes)
- **Adapt to changing conditions** with sliding windows
  - Track non-stationary  $p(t)$  effectively
- **Balance exploration vs exploitation**
  - UCB policies explore efficiently early
  - Exploit optimal code once identified

# Simulation results : Code frequencies for different channel conditions, $\lambda$

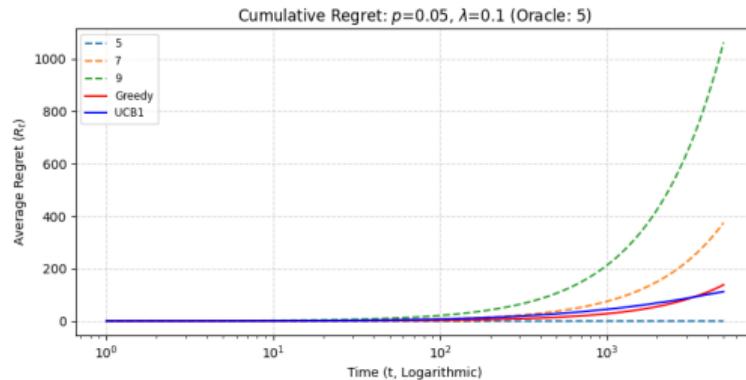
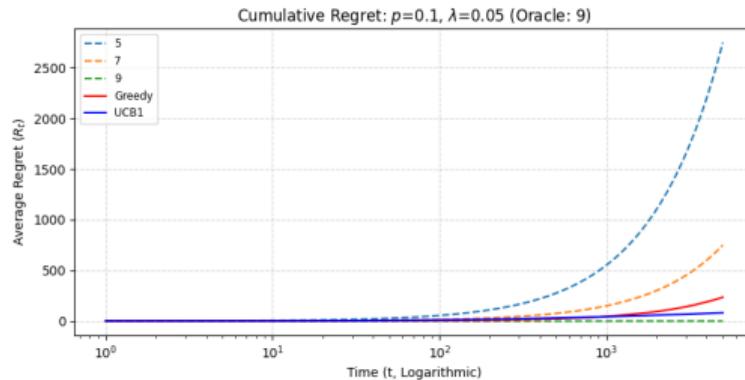
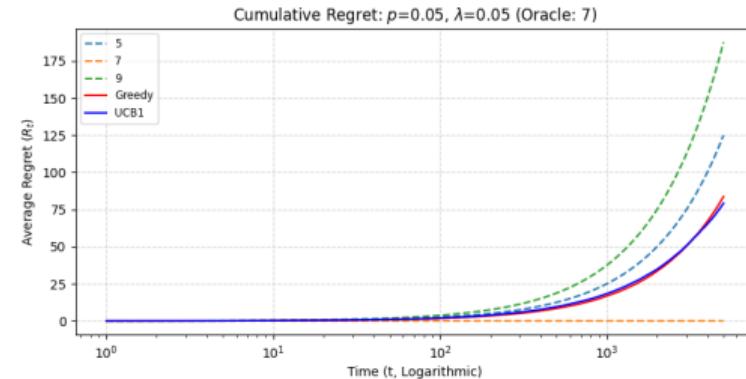
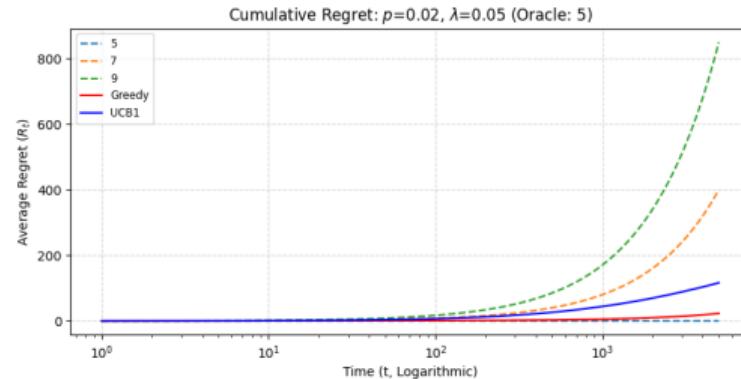


# Simulation results

Table: Cumulative Regret Comparison of Fixed vs. Adaptive QEC Policies ( $T = 5000$  Logical Qubits)

Channel		Optimal Code	Oracle Reward	Fixed Policy Regret			Adaptive Policy Regret	
P	$\lambda$			5-qubit	7-qubit	9-qubit	Greedy	UCB1
0.02	0.05	5-qubit	0.7100	0.00	400.00	850.00	22.85	116.53
0.05	0.05	7-qubit	0.5250	125.00	0.00	187.50	83.67	79.16
0.10	0.05	9-qubit	0.3000	2750.00	750.00	0.00	234.93	81.66
0.05	0.10	5-qubit	0.2500	0.00	375.00	1062.50	138.33	112.22

# Simulation results : Cumulative Regrets



# Simulation results : Total Physical Qubits Transmitted

Table: Total Physical Qubits Transmitted by Fixed and Adaptive Policies ( $T = 5000$  Logical Qubits)

Channel		Optimal Code	Fixed Policy Qubits			Adaptive Policy Qubits	
P	$\lambda$		5-qubit	7-qubit	9-qubit	Greedy	UCB1
0.02	0.05	5-qubit	25000	35000	45000	25568	27837
0.05	0.05	7-qubit	25000	35000	45000	32475	34481
0.10	0.05	9-qubit	25000	35000	45000	42281	44065
0.05	0.10	5-qubit	25000	35000	45000	28589	27624

# Discussion: Why This Matters

## Broader Impact:

- **Resource efficiency:** Use minimum qubits for desired fidelity
  - Critical for NISQ-era devices with limited qubits
- **Robustness:** Automatically adapt to hardware variations
  - No manual recalibration needed
- **Scalability:** Framework extends to more codes
  - Can incorporate surface codes, color codes, etc.
- **Theoretical guarantees:** Provable regret bounds
  - Not just heuristics—mathematically principled

**Future work:** Contextual bandits (incorporate syndrome patterns), multi-objective optimization

# Limitations and Extensions

## Current limitations:

- Assumes i.i.d. noise within windows
- Binary fidelity model (success/failure)
- Fixed code set (3 options)
- Doesn't leverage syndrome structure directly

## Possible extensions:

- **Contextual bandits:** Use syndrome patterns as context
- **Continuous fidelity:** Model full distribution, not just binary
- **Correlated noise:** Model temporal correlations in  $p(t)$
- **Hierarchical codes:** Include concatenated and surface codes
- **Batch decisions:** Choose codes for blocks of qubits jointly

# Conclusion

## Summary:

- Formulated adaptive QEC as a stochastic multi-armed bandit
- Defined reward  $r_t = f_t - \lambda C_{k_t}$  trading fidelity vs cost
- Presented three strategies: Greedy, Parametric UCB, Model-Free UCB
- Connected to our heuristic: greedy policy on estimated  $\hat{p}_t$
- Extended to non-stationary channels via sliding windows

## Key insight:

*Online learning naturally balances exploration of uncertain codes  
with exploitation of codes we believe are optimal*

**Impact:** Principled framework for resource-efficient quantum communication

# Thank You!

**Questions?**

Kollapudi Jithin  
Krunal Vaghela  
Yaswanth R Kumar